

lab6 - openmp

实验环境

系统版本

```
Linux LAPTOP-BGRVTJ4L 5.15.153.1-microsoft-standard-WSL2 #1 SMP Fri Mar 29 23:14:13 UTC
2024 x86_64 x86_64 x86_64 GNU/Linux
```

```
Ubuntu 22.04.4 LTS
```

编译器版本

```
gcc version 11.4.0 (Ubuntu 11.4.0-1ubuntu1~22.04)
```

CPU 物理核数及频率

```
CPU(s):                16
On-line CPU(s) list:    0-15
Thread(s) per core:     2
Core(s) per socket:     8
```

```
CPU: 3792.655MHz
```

四种矩阵乘实现

Naive

Naive gemm 是**最简单**的矩阵乘实现，**并未进行包括分块和多线程在内的优化**，因此也是四种实现中效率最低的实现。

Naive gemm 的 C 矩阵每个元素的计算公式如下：

$$C[i, j] = C[i, j] + A[i, p] \times B[p, j]$$

Naive gemm 通常使用最简单的三重循环实现，下面是其核心代码：

```
/* Macros for row-major order */
#define A(i, j) a[(i) * lda + (j)]
#define B(i, j) b[(i) * ldb + (j)]
#define C(i, j) c[(i) * ldc + (j)]

/* Routine for computing C = A * B + C */
void MY_MMult(int m, int n, int k, double *a, int lda,
              double *b, int ldb,
              double *c, int ldc)
{
```

```

int i, j, p;

for (i = 0; i < m; i++) /* Loop over the rows of C */
{
    for (j = 0; j < n; j++) /* Loop over the columns of C */
    {
        for (p = 0; p < k; p++)
        { /* Update C( i,j ) with the inner product of the ith row of A and the jth
column of B */
            C(i, j) = C(i, j) + A(i, p) * B(p, j);
        }
    }
}

```

Openblas

Openblas 是 **BLAS（基础线性代数程序集）** 的一种开源实现。

使用 openblas 中的 `cblas_dgemm` 函数实现 gemm，核心代码如下：

```

void MY_MMult(int m, int n, int k, double *a, int lda,
              double *b, int ldb,
              double *c, int ldc)
{
    cblas_dgemm(CblasRowMajor, CblasNoTrans, CblasNoTrans, m, n, k,
                1.0, a, lda, b, ldb, 1.0, c, ldc);
}

```

Pthread

POSIX线程（英语：POSIX Threads，常被缩写为 pthreads）是 POSIX 的线程标准，定义了创建和操纵线程的一套 API。实现 POSIX 线程标准的库常被称作 **pthreads**。

使用 pthread 库进行多线程计算，实现 gemm，核心代码如下：

```

#include "defs.h"
#include <pthread.h>
#include <assert.h>
#include <stdio.h>

#define min(a, b) ((a) < (b) ? (a) : (b))
#define max(a, b) ((a) > (b) ? (a) : (b))
#include <math.h>

struct MatrixThreadArgs {
    int m;
    int n;
    int k;
    double *a;
    int lda;
    double *b;
    int ldb;
    double *c;
}

```

```

int ldc;
int section_x_begin;
int section_x_end;
int section_y_begin;
int section_y_end;
};

void *MatrixThreadCalculate(void *arg) {
    struct MatrixThreadArgs matrixThreadArgs = *((struct MatrixThreadArgs *) arg);
    int k = matrixThreadArgs.k;
    double *a = matrixThreadArgs.a;
    int lda = matrixThreadArgs.lda;
    double *b = matrixThreadArgs.b;
    int ldb = matrixThreadArgs.ldb;
    double *c = matrixThreadArgs.c;
    int ldc = matrixThreadArgs.ldc;
    int section_x_begin = matrixThreadArgs.section_x_begin;
    int section_x_end = matrixThreadArgs.section_x_end;
    int section_y_begin = matrixThreadArgs.section_y_begin;
    int section_y_end = matrixThreadArgs.section_y_end;

    const int block_size = min(64, (max(section_x_end - section_x_begin, section_y_end
- section_y_begin)));
    int block_column_num = ceil((section_x_end - section_x_begin) / block_size);
    int block_row_num = ceil((section_y_end - section_y_begin) / block_size);

    for (int block_x = 0; block_x < block_column_num; block_x++) {
        for (int block_y = 0; block_y < block_row_num; block_y++) {
            int block_base_x = section_x_begin + block_x * block_size;
            int block_base_y = section_y_begin + block_y * block_size;
            int block_end_x = min(section_x_end, block_base_x + block_size);
            int block_end_y = min(section_y_end, block_base_y + block_size);

            for (int i = block_base_x; i < block_end_x; i++) {
                for (int j = block_base_y; j < block_end_y; j++) {
                    for (int p = 0; p < k; p++) {
                        C(i, j) = C(i, j) + A(i, p) * B(p, j);
                    }
                }
            }
        }
    }
    return NULL;
}

void MY_MMult(int m, int n, int k, double *a, int lda,
              double *b, int ldb,
              double *c, int ldc) {
    const int x_seperate = 4, y_seperate = 4, thread_num = x_seperate * y_seperate;

    pthread_t threads[thread_num];
    struct MatrixThreadArgs sectionMatrixThreadArgs[thread_num];
    int rc;

    int section_weight = ceil(n / x_seperate);
    int section_height = ceil(m / y_seperate);

```

```

int thread_index = 0;
for (int section_x = 0; section_x < x_separate; section_x++) {
    for (int section_y = 0; section_y < y_separate; section_y++) {
        int section_x_begin = section_x * section_weight;
        int section_x_end = min(n, section_x_begin + section_weight);
        int section_y_begin = section_y * section_height;
        int section_y_end = min(m, section_y_begin + section_height);

        sectionMatrixThreadArgs[thread_index] = (struct MatrixThreadArgs){
            .m = m,
            .n = n,
            .k = k,
            .a = a,
            .lda = lda,
            .b = b,
            .ldb = ldb,
            .c = c,
            .ldc = ldc,
            .section_x_begin = section_x_begin,
            .section_x_end = section_x_end,
            .section_y_begin = section_y_begin,
            .section_y_end = section_y_end
        };
        rc = pthread_create(&threads[thread_index], NULL, MatrixThreadCalculate,
&sectionMatrixThreadArgs[thread_index]);
        assert(rc == 0);
        thread_index++;
    }
}
for (int i = 0; i < thread_num; i++) {
    rc = pthread_join(threads[i], NULL);
    assert(rc == 0);
}
}

```

Openmp

OpenMP (Open Multi-Processing) 是一套支持跨平台共享内存方式的多线程并发的编程 API。

核心代码如下：

```

void MY_MMult(int m, int n, int k, double *a, int lda,
              double *b, int ldb,
              double *c, int ldc)
{
    int i, j, p;
    #pragma omp parallel for private(j, p)
    for (i = 0; i < m; i++) /* Loop over the rows of C */
    {
        for (j = 0; j < n; j++) /* Loop over the columns of C */
        {
            for (p = 0; p < k; p++)
            { /* Update C( i,j ) with the inner product of the ith row of A and the jth
column of B */

```


```

    C(i, j) = C(i, j) + A(i, p) * B(p, j);
    // printf("Thread %d: i=%d, j=%d, p=%d\n", omp_get_thread_num(), i, j, p);
}
}
// printf("Thread %d: i=%d\n", omp_get_thread_num(), i);
}
}

```

Gflops

下图为四种 gemm 实现在不同矩阵规模下的 gflops 曲线图：



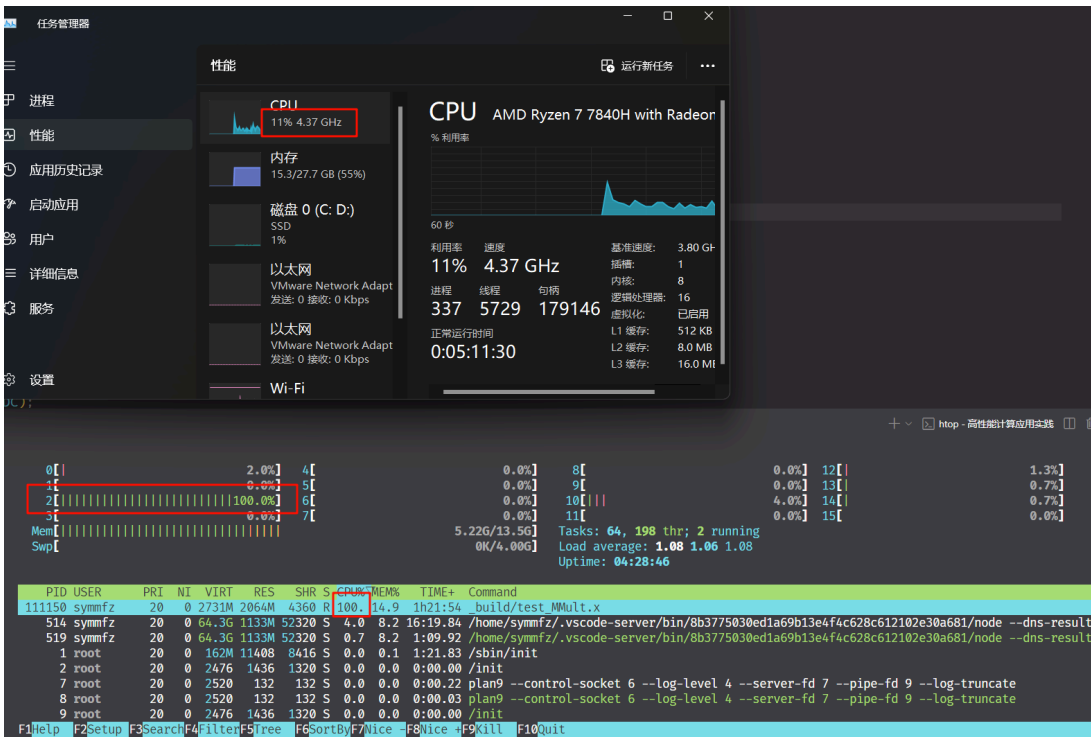
根据上图，不难发现以下结论：

- Cblas 实现的 gemm 在绝大多数矩阵规模下，gflops 都显著高于其他四种实现，最高可达 naive 实现的 500 倍以上
- Cblas 实现在较低矩阵规模时 gflops 较低，随着矩阵规模的上涨 gflops 值先快速上涨然后稳定。Gflops 峰值出现在 1024×1024 的矩阵规模附近，峰值大小约为 64.35
- Openmp 和 pthread 实现的 gflops 曲面相似，从 gflops 的大小上看，大于 naive 实现并显著小于 cblas 实现；从曲线的变化上看，随着矩阵规模的增大，gflops 先增后减，峰值出现在 512×512 或 1024×1024 附近，峰值 gflops 约为 4.4
- Naive 实现的 gflops 值在矩阵规模大于 32×32 时最低。从趋势上看，gflops 值大致随着矩阵规模的增大而减小

截图

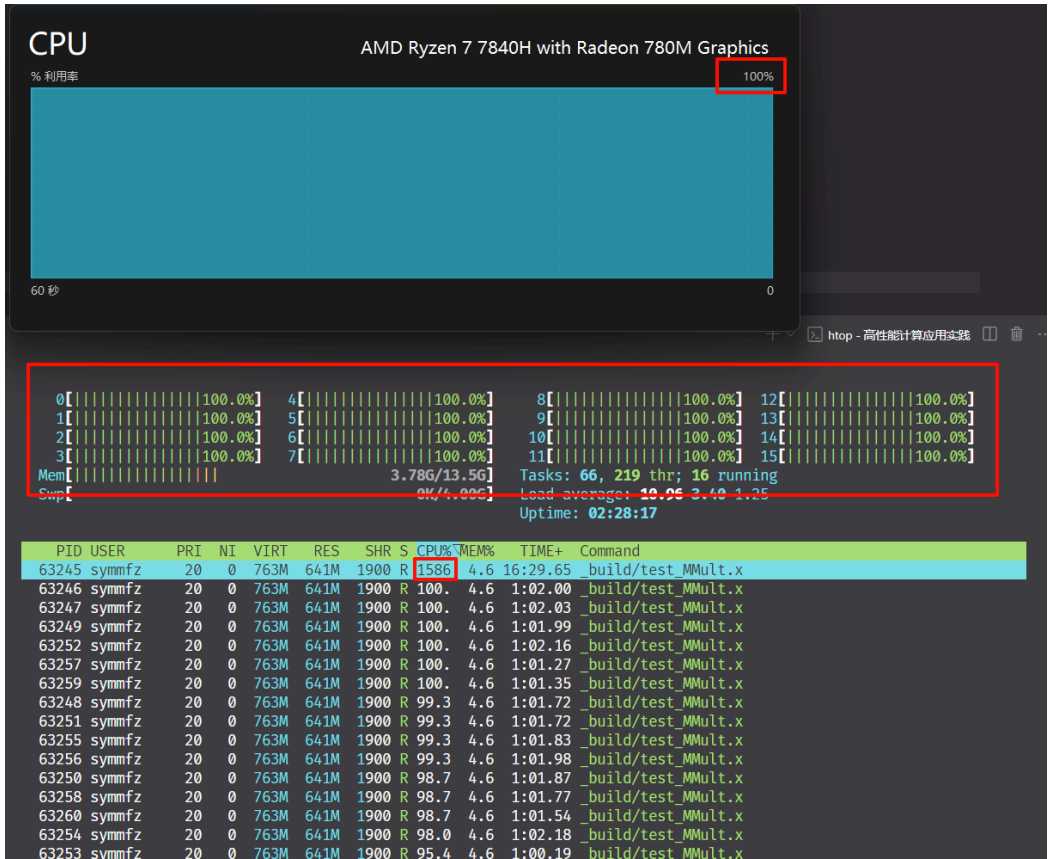
Naive

Naive 为单线程运行，无法充分利用 CPU 的性能。




openblas

Openblas 可以充分利用 CPU 性能。




Pthread

Pthread 16 线程运行



Openmp

Openmp 使用 16 线程运行，可以充分利用 CPU 性能。



PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
88539	symmfb	20	0	289908	166092	1920	R	1582	1.2	3:30.61	test_MMult.x

```

fish—sh—sh—sh—node—node—fish—test_MMult.x—15*[{test_MMult.x}]
                        |
                        |—fish—pstree
                        |   |
                        |   |—2*[{fish}]
                        |   |
                        |   |—12*[{node}]
                        |   |
                        |   |—node—12*[{node}]
                        |   |
                        |   |—node—cpptools—25*[{cpptools}]
                        |   |   |
                        |   |   |—node—10*[{node}]
                        |   |   |   |
                        |   |   |   |—2*[{node—6*[{node}]}]
                        |   |   |   |
                        |   |   |   |—16*[{node}]
                        |   |   |
                        |   |   |—10*[{node}]
                        |   |
                        |   |—10*[{node}]

```

上面这行代码表示将程序的输出追加到性能数据文件中。

Lab 5 - thread

截图

```
top - 00:44:48 up 1:09, 1 user, load average: 2.24, 1.59, 1.11
Tasks: 1 total, 0 running, 1 sleeping, 0 stopped, 0 zombie
%Cpu(s): 52.3 us, 0.6 sy, 0.0 ni, 45.7 id, 0.0 wa, 0.0 hi, 1.4 si, 0.0 st
MiB Mem : 13824.3 total, 11132.9 free, 2135.1 used, 556.3 buff/cache
MiB Swap: 4096.0 total, 4096.0 free, 0.0 used. 11426.7 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
37461	symmfz	20	0	232364	165904	1776	S	800.0	1.2	3:11.61	test_MMult.x

```
top - 00:40:05 up 1:04, 1 user, load average: 2.72, 1.56, 0.97
Threads: 9 total, 8 running, 1 sleeping, 0 stopped, 0 zombie
%Cpu(s): 49.6 us, 0.3 sy, 0.0 ni, 48.8 id, 0.0 wa, 0.0 hi, 1.2 si, 0.0 st
MiB Mem : 13824.3 total, 11130.1 free, 2138.7 used, 555.5 buff/cache
MiB Swap: 4096.0 total, 4096.0 free, 0.0 used. 11423.2 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
35215	symmfz	20	0	232364	165884	1760	R	99.9	1.2	0:13.17	test_MMult.x
35216	symmfz	20	0	232364	165884	1760	R	99.9	1.2	0:13.17	test_MMult.x
35217	symmfz	20	0	232364	165884	1760	R	99.9	1.2	0:13.17	test_MMult.x
35218	symmfz	20	0	232364	165884	1760	R	99.9	1.2	0:13.17	test_MMult.x
35219	symmfz	20	0	232364	165884	1760	R	99.9	1.2	0:13.17	test_MMult.x
35220	symmfz	20	0	232364	165884	1760	R	99.9	1.2	0:13.17	test_MMult.x
35221	symmfz	20	0	232364	165884	1760	R	99.9	1.2	0:13.17	test_MMult.x
35222	symmfz	20	0	232364	165884	1760	R	99.9	1.2	0:13.17	test_MMult.x
34099	symmfz	20	0	232364	165884	1760	S	0.0	1.2	1:39.55	test_MMult.x